

■ Duoluga Learning Pipeline

This document explains how Duoluga learns in two stages: 1. ****Short-term (RAG):**** Data is tagged and stored in memory, then used to augment prompts so the LLM can answer in Luo/Swahili/Kikuyu immediately. 2. ****Long-term (Fine-tuning):**** The same dataset is later used to fine-tune the base LLM (Qwen now, BitNet later), so the model permanently learns the languages and cultural context. The diagram (Duoluga_Learning_Pipeline.png) should be saved alongside this PDF in your project folder.